



A.D. 1308

unipg

DIPARTIMENTO
DI INGEGNERIA

UNIVERSITÀ DEGLI STUDI DI PERUGIA

Dipartimento di Ingegneria

Laurea Magistrale in
Ingegneria Informatica e Robotica

**Un fantastico titolo
per la mia tesi di laurea!**

Relatore:

Prof. Gabriele Costante

Candidato:

Paolo Speziali

Anno Accademico 2023/2024

Contents

1	Introduction	7
2	State of the Art	9
3	Prior Knowledge	11
3.1	Deep Learning	12
3.2	Reinforcement Learning	14
3.3	Causal Inference	18
3.4	Generative Adversarial Networks	18
4	Methodology	19
5	Experiments	21
6	Conclusion	23

List of Figures

3.1	A feed-forward neural network with two hidden layers	12
3.2	A simple reinforcement learning task	14
3.3	The example grid as a Markov Process graph, where the nodes represent the states and the edges represent the state transitions.	17

Chapter 1

Introduction

Chapter 2

State of the Art

Chapter 3

Prior Knowledge

3.1 Deep Learning

In the field of machine learning, the first models you are often introduced to are those for regression and classification that utilize linear combinations of fixed basis functions. According to [Bishop, 2008], these models possess useful analytical and computational properties, but their practical applicability is constrained since their capacity is limited to linear functions and they cannot understand the interaction between any two input variables. This leads to problems such as the curse of dimensionality, where the number of possible interactions between variables grows exponentially with the number of input variables. To address large-scale problems, it is essential to adapt the basis functions to the data, as demonstrated by support vector machines (SVMs). Alternatively, one can fix the number of basis functions in advance while allowing them to be adaptive, utilizing parametric forms for the basis functions with parameter values adjusted during training.

The most successful example of this approach in pattern recognition is the **feed-forward neural network**, also known as the multilayer perceptron (MLP).

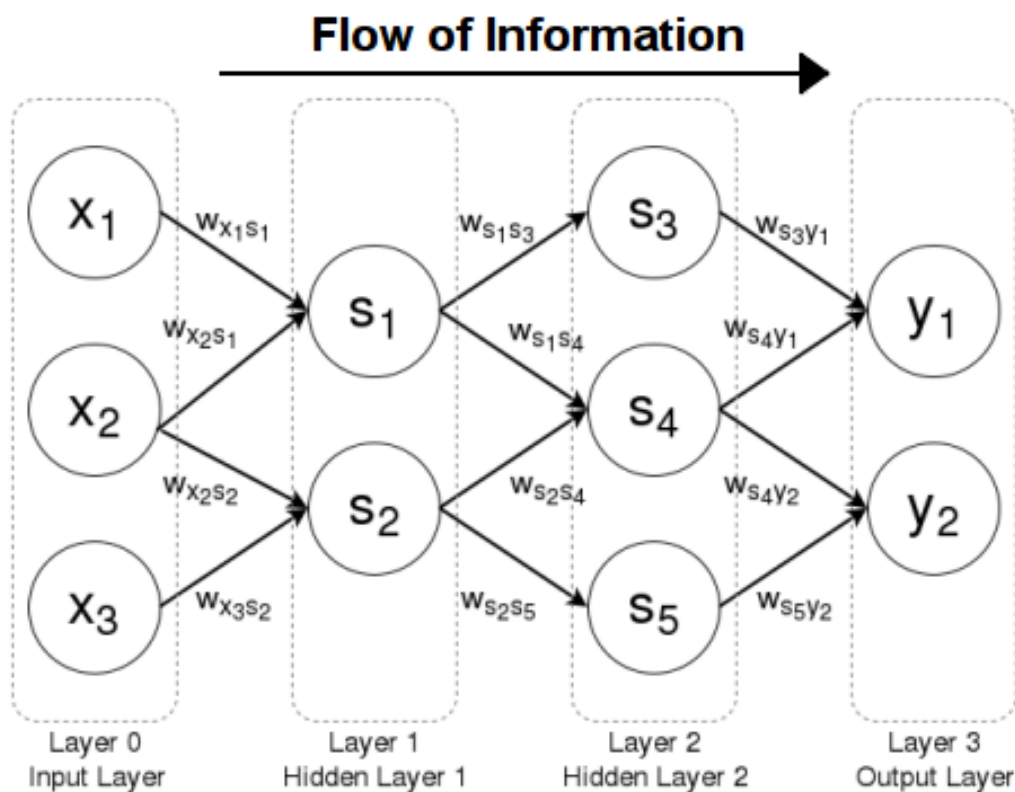


Figure 3.1: A feed-forward neural network with two hidden layers

Source: Brilliant

As explained in [Goodfellow et al., 2016], the goal of a feedforward network is to approximate a specific function f^* . For example, in a classification task, $y = f^*(x)$ maps an input x to a category y . A feedforward network defines a mapping

$y = f(x; \theta)$ and learns the parameters θ that yield the best approximation of this function.

These models are called **feedforward** because information flows unidirectionally from the input x , through intermediate computations defining f , and ultimately to the output y , without feedback loops.

The training data directly specifies the output layer's behavior but not the intermediate layers, which are called **hidden layers** because their desired outputs are not provided by the training data. The learning algorithm must determine how to best utilize these hidden layers to approximate f^* . Every layer of the network computes a non-linear transformation of the previous layer's activations, this way a complex function can be approximated by composing simpler functions, one for each layer. Layers are composed of a set of units, where each unit is a node that computes a non-linear function of the weighted sum of its inputs and is only connected to units in the previous and the following layer.

A feedforward network is called a **deep neural network** if it has more than one hidden layer, the branch of machine learning that studies deep neural networks is called **deep learning**.

They are called networks due to their structure, which involves composing multiple functions together. Typically, this composition is represented by a directed acyclic graph. For instance, functions $f^{(1)}$, $f^{(2)}$, and $f^{(3)}$ might be connected in a chain to form $f(x) = f^{(3)}(f^{(2)}(f^{(1)}(x)))$.

Functions $f^{(1)}$ and $f^{(2)}$ must be non linear, otherwise the composition would collapse into a single linear function. These non-linear functions are called **activation functions**. The last function $f^{(3)}$ is typically a linear function that maps the output of the final hidden layer to the output layer. This is the same as applying a linear model to a transformed input $\phi(x)$, where ϕ is a nonlinear transformation. The question then becomes how to choose the mapping ϕ .

The strategy of deep learning is to learn ϕ : in this approach, we use a model

$$y = f(x; \theta, w) = \phi(x; \theta)^\top w$$

Here, we have parameters θ that are used to learn ϕ from a broad class of functions, and parameters w that map $\phi(x)$ to the desired output. This approach allows for greater flexibility: specifically, by using a broad family of functions $\phi(x; \theta)$, the human designer only needs to select the appropriate general function family rather than finding precisely an exact function.

The **universal approximation theorem** [Hornik et al., 1989] tells us that regardless of what function we are trying to learn, we know that a feedforward network with enough units will be able to *represent* this function, however we are not guaranteed that the training algorithm will be able to *learn* it.

3.2 Reinforcement Learning

“**Reinforcement learning** is learning [...] how to map situations to actions so as to maximize a numerical reward signal. The learner is not told which actions to take, but instead must discover which actions yield the most reward by trying them. In the most interesting and challenging cases, actions may affect not only the immediate reward but also the next situation and, through that, all subsequent rewards. These two characteristics, **trial-and-error search** and **delayed reward**, are the two most important distinguishing features of reinforcement learning.” [Sutton and Barto, 1998]

Let’s explore the basic concepts of reinforcement learning with a simple example as described in [Zhao, 2024]. Consider a grid world scenario as shown in Figure 3.2 where a robot, referred to as the **agent**, moves between cells, occupying one cell at a time. The white cells are accessible, while the orange cells are forbidden. The goal is for the agent to reach a target cell.

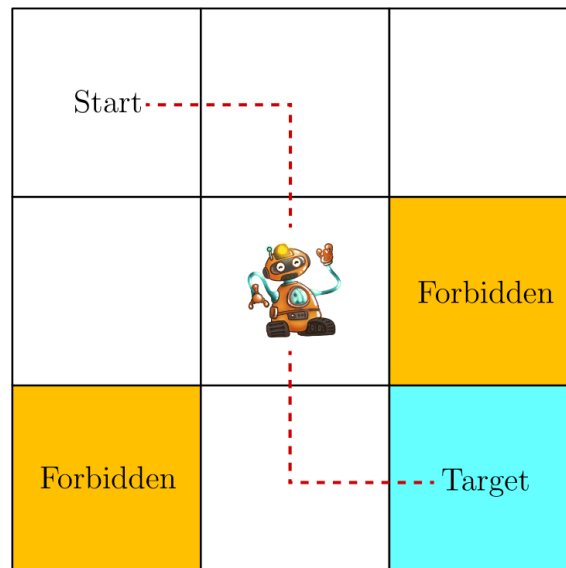


Figure 3.2: A simple reinforcement learning task

Source: [Zhao, 2024]

To achieve this, the agent needs a **policy** that guides it to the target efficiently, avoiding forbidden cells and unnecessary detours. If the grid layout is known, planning is simple. However, without prior knowledge, the agent must explore and learn through trial and error.

The agent’s status in the grid is defined by its **state** $s_i \in \mathcal{S}$, which represents its location relative to the **environment**. In the examples with nine cells, the state space is $\mathcal{S} = \{s_1, s_2, \dots, s_9\}$.

From each state, the agent can perform five **actions**: move up, right, down, left or stay put, denoted as a_1, a_2, \dots, a_5 , this set of actions is the action space $\mathcal{A} = \{a_1, \dots, a_5\}$. The available actions can vary by state, for instance, at s_1 ,

actions a_1 (up) and a_4 (left) would collide with the grid boundary, so the action space is $\mathcal{A}(s_1) = \{a_2, a_3, a_5\}$.

When taking an action, the agent may move from one state to another, a process known as **state transition**. For example, if the agent is at state s_1 and selects action a_2 (moving rightward), it transitions to state s_2 . This process can be represented as:

$$s_1 \xrightarrow{a_2} s_2$$

The state transition process is defined for each state and its associated actions. Mathematically, state transitions are described by conditional probabilities. For example, for s_1 and a_2 , the conditional probability distribution is:

$$\begin{aligned} p(s_1 \mid s_1, a_2) &= 0, \\ p(s_2 \mid s_1, a_2) &= 1, \\ p(s_3 \mid s_1, a_2) &= 0, \\ p(s_4 \mid s_1, a_2) &= 0, \\ p(s_5 \mid s_1, a_2) &= 0, \end{aligned}$$

indicating that taking a_2 at s_1 guarantees the agent moves to s_2 , with a probability of one, and zero probability for other states.

This is a deterministic state transition, but state transitions can also be stochastic, requiring conditional probability distributions. For instance, if random wind gusts affect the grid, taking action a_2 at s_1 might blow the agent to s_5 instead of s_2 , resulting in $p(s_5 \mid s_1, a_2) > 0$.

A **policy** is a function that maps states to actions, indicating the agent's behavior in the environment. In other words, a policy tells the agent what action to take at each state.

Mathematically, policies can be described by conditional probabilities. Denote the policy in Figure 1.4 as $\pi(a \mid s)$, a conditional probability distribution function defined for every state. For example, the policy for s_1 is:

$$\begin{aligned} \pi(a_1 \mid s_1) &= 0, \\ \pi(a_2 \mid s_1) &= 1, \\ \pi(a_3 \mid s_1) &= 0, \\ \pi(a_4 \mid s_1) &= 0, \\ \pi(a_5 \mid s_1) &= 0, \end{aligned}$$

indicating that the probability of taking action a_2 at state s_1 is one, and the probabilities of taking other actions are zero.

The above policy is deterministic, but policies may generally be stochastic. For instance, the policy in Figure 1.5 is stochastic: at state s_1 , the agent may take actions to move either rightward or downward, each with a probability of 0.5. In

this case, the policy for s_1 is:

$$\begin{aligned}\pi(a_1 | s_1) &= 0, \\ \pi(a_2 | s_1) &= 0.5, \\ \pi(a_3 | s_1) &= 0.5, \\ \pi(a_4 | s_1) &= 0, \\ \pi(a_5 | s_1) &= 0.\end{aligned}$$

After executing an action at a state, the agent receives a reward r as feedback from the environment. The **reward** is a function of the state s and action a , denoted as $r(s, a)$ and it can be positive, negative, or zero. Different rewards influence the policy the agent learns. Generally, a positive reward encourages the agent to take the corresponding action, while a negative reward discourages it.

However we can't find good policies by simply selecting actions with the greatest immediate rewards since they do not consider long-term outcomes. To determine a good policy, we must consider the total reward obtained in the long run and an action with the highest immediate reward may not lead to the greatest total reward.

A **trajectory** is a state-action-reward chain. For example, the agent in our example may follow this trajectory:

$$s_1 \xrightarrow{a_2, r=0} s_2 \xrightarrow{a_3, r=0} s_5 \xrightarrow{a_3, r=0} s_8 \xrightarrow{a_2, r=1} s_9.$$

The return of this trajectory is the sum of all rewards collected along it:

$$\text{return} = 0 + 0 + 0 + 1 = 1$$

Returns, also called total rewards or cumulative rewards, are used to evaluate policies. Returns can also be defined for infinitely long trajectories, which may diverge. Therefore, we introduce the concept of **discounted return** for infinitely long trajectories. The discounted return is the sum of the discounted rewards:

$$\text{discounted return} = 0 + \gamma \cdot 0 + \gamma^2 \cdot 0 + \gamma^3 \cdot 1 + \gamma^4 \cdot 1 + \gamma^5 \cdot 1 + \dots,$$

where $\gamma \in (0, 1)$ is the discount rate. If γ is close to 0, the agent emphasizes near-future rewards, resulting in a short-sighted policy. If γ is close to 1, the agent emphasizes far-future rewards.

When following a policy, the agent may stop at **terminal states**, resulting in a trajectory called an **episode**.

The framework of **Markov Decision Processes** (MDPs) allows us to formally presents these concepts. An MDP is a general framework for describing stochastic dynamical systems and its key components are:

- **Sets:**

- **State space:** The set of all states, denoted as \mathcal{S} .

- **Action space:** The set of actions, denoted as $\mathcal{A}(s)$, associated with each state $s \in \mathcal{S}$.
- **Reward set:** The set of rewards, denoted as $\mathcal{R}(s, a)$, associated with each state-action pair (s, a) .
- **Model:**
 - **State transition probability:** At state s , when taking action a , the probability of transitioning to state s' is $p(s' | s, a)$.
It holds that $\sum_{s' \in \mathcal{S}} p(s' | s, a) = 1 \quad \forall (s, a)$.
 - **Reward probability:** At state s , when taking action a , the probability of obtaining reward r is $p(r | s, a)$.
It holds that $\sum_{r \in \mathcal{R}(s, a)} p(r | s, a) = 1 \quad \forall (s, a)$.
- **Policy:** At state s , the probability of choosing action a is $\pi(a | s)$.
It holds that $\sum_{a \in \mathcal{A}(s)} \pi(a | s) = 1 \quad \forall s \in \mathcal{S}$.
- **Markov property:** The Markov property refers to the **memoryless** property of a stochastic process. Mathematically, it means that

$$p(s_{t+1} | s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = p(s_{t+1} | s_t, a_t),$$

$$p(r_{t+1} | s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = p(r_{t+1} | s_t, a_t),$$

where t represents the current time step and $t+1$ represents the next time step. This indicates that the next state or reward depends only on the current state and action and is independent of the previous ones.

When the policy in a MDP is fixed, it reduces to a **Markov Process** (MP), this transformation simplifies the MDP by eliminating the decision-making aspect. A Markov process is referred to as a **Markov Chain** if it operates in discrete time and the number of states is either finite or countable.

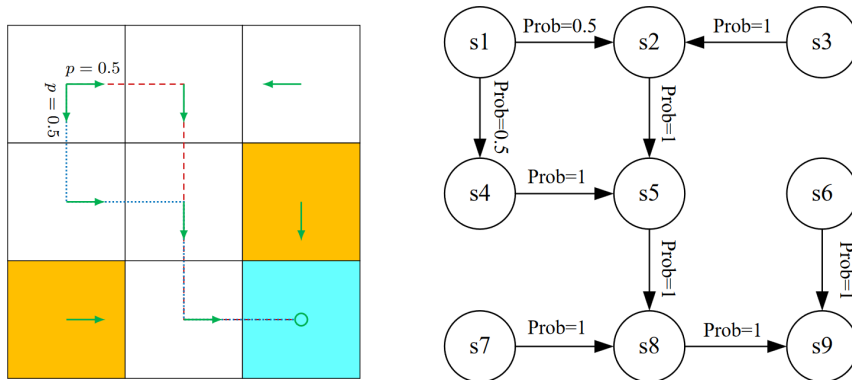


Figure 3.3: The example grid as a Markov Process graph, where the nodes represent the states and the edges represent the state transitions.

Source: [Zhao, 2024]

3.3 Causal Inference

[Neal, 2020]

3.4 Generative Adversarial Networks

Chapter 4

Methodology

Chapter 5

Experiments

Chapter 6

Conclusion

Bibliography

- [Bishop, 2008] Bishop, C. M. (2008). *Pattern Recognition and Machine Learning*. Springer-Verlag, New York.
- [Goodfellow et al., 2016] Goodfellow, I., Bengio, Y., and Courville, A. (2016). *Deep Learning*. Adaptive computation and machine learning. MIT Press.
- [Hornik et al., 1989] Hornik, K., Stinchcombe, M., and White, H. (1989). Multilayer feedforward networks are universal approximators. *Neural Networks*, 2(5):359–366.
- [Neal, 2020] Neal, B. (2020). Introduction to causal inference.
- [Sutton and Barto, 1998] Sutton, R. S. and Barto, A. G. (1998). *Reinforcement Learning: An Introduction*. The MIT Press.
- [Zhao, 2024] Zhao, S. (2024). *Mathematical Foundations of Reinforcement Learning*. Springer Nature Press.